

Graph Neural Networks (GNNs)

Introduction, Concepts, and Applications

Ali Balaj

FZ Business Informatics

December 9, 2024

Table of Contents

- 1 Introduction to GNNs
- 2 Core Concepts of GNNs
- 3 Applications of GNNs
- 4 Challenges and Future Directions

What are Graph Neural Networks?

- **Definition:** GNNs are neural networks designed to operate on graph-structured data, leveraging nodes, edges, and relationships.
- **Purpose:** They encode graph structures into feature-rich representations for tasks like node classification, link prediction, and graph classification.
- **Example Applications:**
 - Social Network Analysis
 - Recommendation Systems
 - Molecular Graphs in Drug Discovery

Why Graph Neural Networks?

- Traditional neural networks are not suited for graph data.
- GNNs leverage the rich connectivity and structural information inherent in graphs.
- Key advantages:
 - Capturing complex relationships between entities.
 - Directly encoding graph structures into embeddings.

- **Graph Components:**

- **Nodes (Vertices):** Entities in the graph (e.g., users, molecules).
- **Edges:** Relationships or connections between nodes.
- **Node Features:** Attributes of nodes (e.g., user preferences, atom types).
- **Edge Features:** Attributes of edges (e.g., weights, interaction types).

- **Mathematical Representation:**

- A graph $G = (V, E)$, where V is the set of nodes and E is the set of edges.
- Node features: $X \in \mathbb{R}^{|V| \times d}$, where d is the feature dimension.
- Adjacency matrix: $A \in \mathbb{R}^{|V| \times |V|}$.

Key Operations in GNNs

- **Message Passing:**

- Nodes exchange information with their neighbors.
- Messages are aggregated to update node embeddings.

- **Aggregation Function:**

- Combines messages from neighbors (e.g., summation, mean, max).
- Ensures permutation invariance.

- **Update Function:**

- Updates node embeddings based on aggregated messages.
- Often implemented as a neural network layer.

Graph Convolutional Network (GCN)

- **GCN Layer:**

$$H^{(l+1)} = \sigma \left(\hat{A} H^{(l)} W^{(l)} \right)$$

where:

- \hat{A} : Normalized adjacency matrix.
 - $H^{(l)}$: Node embeddings at layer l .
 - $W^{(l)}$: Trainable weights.
 - σ : Activation function (e.g., ReLU).
- **Intuition:** Each node aggregates information from its neighbors and updates its embedding.

- **Node Classification:**

- Predict node labels based on graph structure and features.
- Example: Predict user behavior in a social network.

- **Link Prediction:**

- Predict missing or future connections in a graph.
- Example: Recommending friends on a social platform.

- **Graph Classification:**

- Classify entire graphs.
- Example: Classify molecules based on their properties.

Challenges of GNNs

- Scalability: Training GNNs on large graphs is computationally expensive.
- Over-smoothing: Node embeddings may become indistinguishable in deep GNNs.
- Dynamic Graphs: Adapting GNNs to handle evolving graph structures.

- Designing scalable architectures for large-scale graphs.
- Enhancing interpretability of GNN predictions.
- Exploring hybrid models combining GNNs with other machine learning techniques.

Thank you!
Any questions?